

Implementation and Performance Evaluation of the QUIC Protocol in Linux Kernel

Peng Wang, Carmine Bianco, Janne Riihijärvi and Marina Petrova

Presenter: Kota Yatagai

Abstract

- QUIC is a UDP-based transport protocol designed to improve HTTPS performance, overcoming problems in TCP
- Current QUIC deployments run in user space, but context switching per packet limits its performance
- This study implements QUIC in the Linux kernel to fairly compare it against TCP
- Experiments show that kernel-space QUIC outperforms TCP in scenarios with low latency and high packet loss

Introduction: TCP and its problems

- TCP, designed in 1974, has evolved but struggles with modern requirements
 - TCP has no built-in support for encryption
 - TLS over TCP (HTTP/2) improved security but added handshake delays
 - TCP's strict reliable streams limits performance in real-time communications
- Updating TCP is impractical because of the “protocol ossification”
 - it would require global update to OS kernels and middleboxes

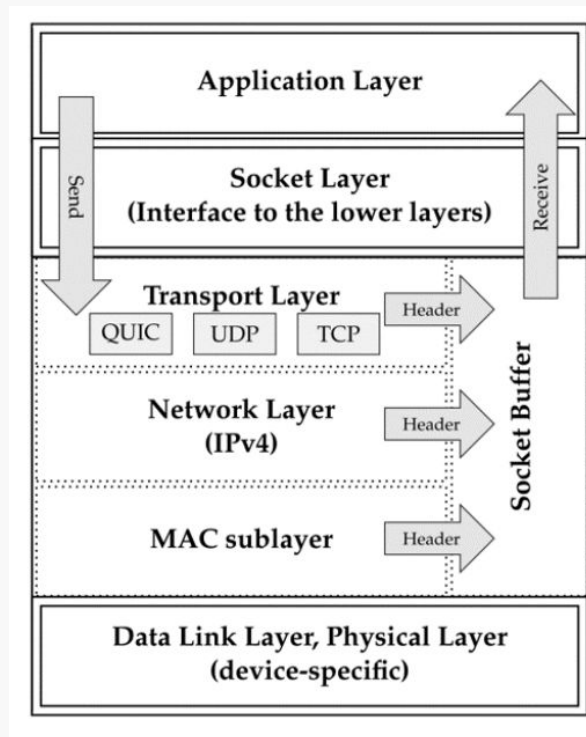
Introduction: Rise of QUIC and unfair comparison

- QUIC is designed to run in user space, making deployment quicker
- QUIC also has TLS1.3 as its required subsystem, ensuring security
 - QUIC-specific information such as initial CID is piggybacked in TLS handshake
 - Whole QUIC handshake will be done with 1RTT at maximum
- Prior studies compared user-space QUIC vs kernel-space TCP, which seemed unfair
- This work implements QUIC inside the Linux kernel for a fair comparison against TCP

Implementation: Protocol registration

QUIC initialization in kernel

- Modify kernel boot process (inet_init) so following steps are executed
 - Register QUIC's socket operations
 - Assign IPPROTO_QUIC
 - Map SOCK_DGRAM and QUIC



Implementation: Connection establishment

- a QUIC connection is established with the berkeley Socket API
 - socket(), bind(), then connect() or listen(), and then accept()

TLS handshake is not implemented in this paper

- Thus following evaluation has to be perceived as TCP vs “QUIC plain text”
 - which is, somehow fair? but definitely impractical


Evaluation: Virtual testbed setup

- Environment: Oracle VirtualBox
 - Both the client and server run on separate VMs
 - Single-core isolated CPU (100% host CPU usage allowed) with 3GB RAM
 - Intel PRO/1000 MT Gigabit Ethernet adaptor
- Host machine: Intel Core i7-6700 (8 cores, 16GB RAM)
- OS: Ubuntu 14.04 with Linux kernel 3.13.11
- Network: emulated with netem (tc)
- Method: 20 repetition of file transfer of size 2.46MB


Evaluation: Wireless testbed setup

- Environment: Up to 3 PCs with (probably) the same CPU, RAM and OS
 - One for the server and one or two for the client depending on the measurement
- Network: IEEE 802.11 (WiFi) with the ASUS RT-AC66U router
- Method: Unknown repetition of file transfer of size 39.69MB


Virtual testbed measurement: Throughput in multiple packet loss rates



(a). No packet loss



(b) 0.1% packet loss rate

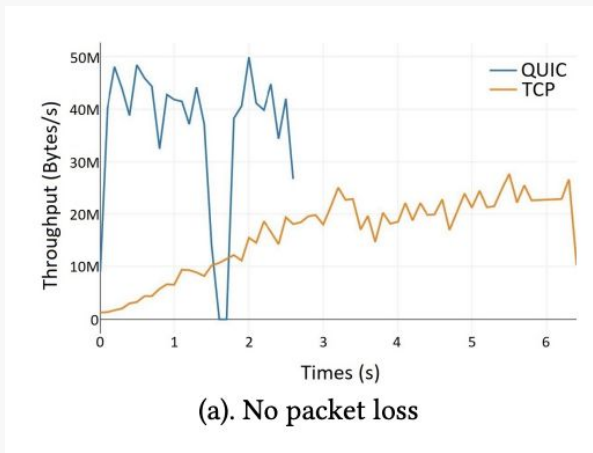


(c) 1% packet loss rate

Figure 3: Throughput of QUIC and TCP with different packet loss rate with 10 ms RTT.

- Repetitions of a single file transfer with 0/0.1/1% packet loss rate
- RTT is fixed to 10ms
- QUIC always outperforms TCP but may have a higher jitter for the low loss rate cases

Out-of-paper topic: Why high throughput and high jitter?




- This paper's implementation does not limit data size of a single transfer
 - as long as cwnd allows, the kernel sends queued packets all at once (bulk)
 - RFC9000 limits this with MAX_DATA or MAX_STREAM_DATA
- Basically not having flow control makes high throughput and high jitter


Out-of-paper topic: What if the implementation is proper?

- Theoretically, QUIC will not overcome TCP in terms of throughput
 - because encryption process is computational intensive.
- QUIC will overcome TCP + TLS1.3
 - quicker handshakes, more flexible SACK (Selective ACK)
- We do not have a real-world measurement of properly implemented QUIC in kernel at this point
 - there are other in-kernel QUIC implementations such as lxin/quic, though not evaluated yet


Virtual testbed measurement: multiple RTTs with multiple packet loss rates



Packet loss rate
(a) 10 ms RTT.



Packet loss rate
(b) 30 ms RTT.



Packet loss rate
(c) 100 ms RTT.

- QUIC always outperforms TCP
 - better loss recovery systems
- The performance deteriorates by roughly a factor of 10 for both QUIC and TCP when just 0.1% loss is introduced into the network

out-of-paper topic: QUIC vs TCP in loss recovery



- Linux TCP employs DUPACK and SACK
 - RTO(Retransmission Timeout) is slow (usually 200ms~)
 - Trigger retransmission with 3 duplicate ACKs (DUPACK)
 - also shrink cwnd for congestion avoidance
 - normal ACKs only notify a single packet number...
 - Selective ACK notifies which packets are missing
 - Less insufficient retransmission!

out-of-paper topic: QUIC vs TCP in loss recovery


- QUIC employs PTO and ACK Ranges
 - RTO is slow
 - Set Probe TimeOut (usually 1.125 RTT) when sending packets
 - If an ACK doesn't come back by PTO, retransmit
 - PTO increases exponentially until the maximum timeout
 - Faster than the TCP's DUPACK (PTO vs waiting for 3 ACKs)
 - QUIC also has a better version of SACK (no limit for packet indication)
 - even with huge loss, QUIC is expected to recover faster

Virtual testbed: Link sharing


- Running QUIC and TCP in two virtual machines
- the bandwidth is fixed to 2MBps
- the link is shared fairly between the two because both use the same CCA



Wireless testbed: 0% loss rate




(a) 10 ms RTT and 0% loss rate.




(c) 50 ms RTT and 0% loss rate.

- In a normal network situation (0 packet loss), QUIC reaches its peak throughput quicker than TCP
 - quicker handshake, even including TLS

Wireless testbed: 10% loss rate



(b) 10 ms RTT and 10% loss rate.



(c) 100 ms RTT and 10% loss rate.

- QUIC has higher throughput over TCP
 - better loss recovery (faster retransmission, faster recovery)

Conclusion

- QUIC outperforms TCP in lossy networks
- QUIC also completes handshake quicker
- In-kernel QUIC with no encryption has higher performance in most situations
 - proper QUIC vs TCP+TLS1.3 will be the same result, but there's no measurement